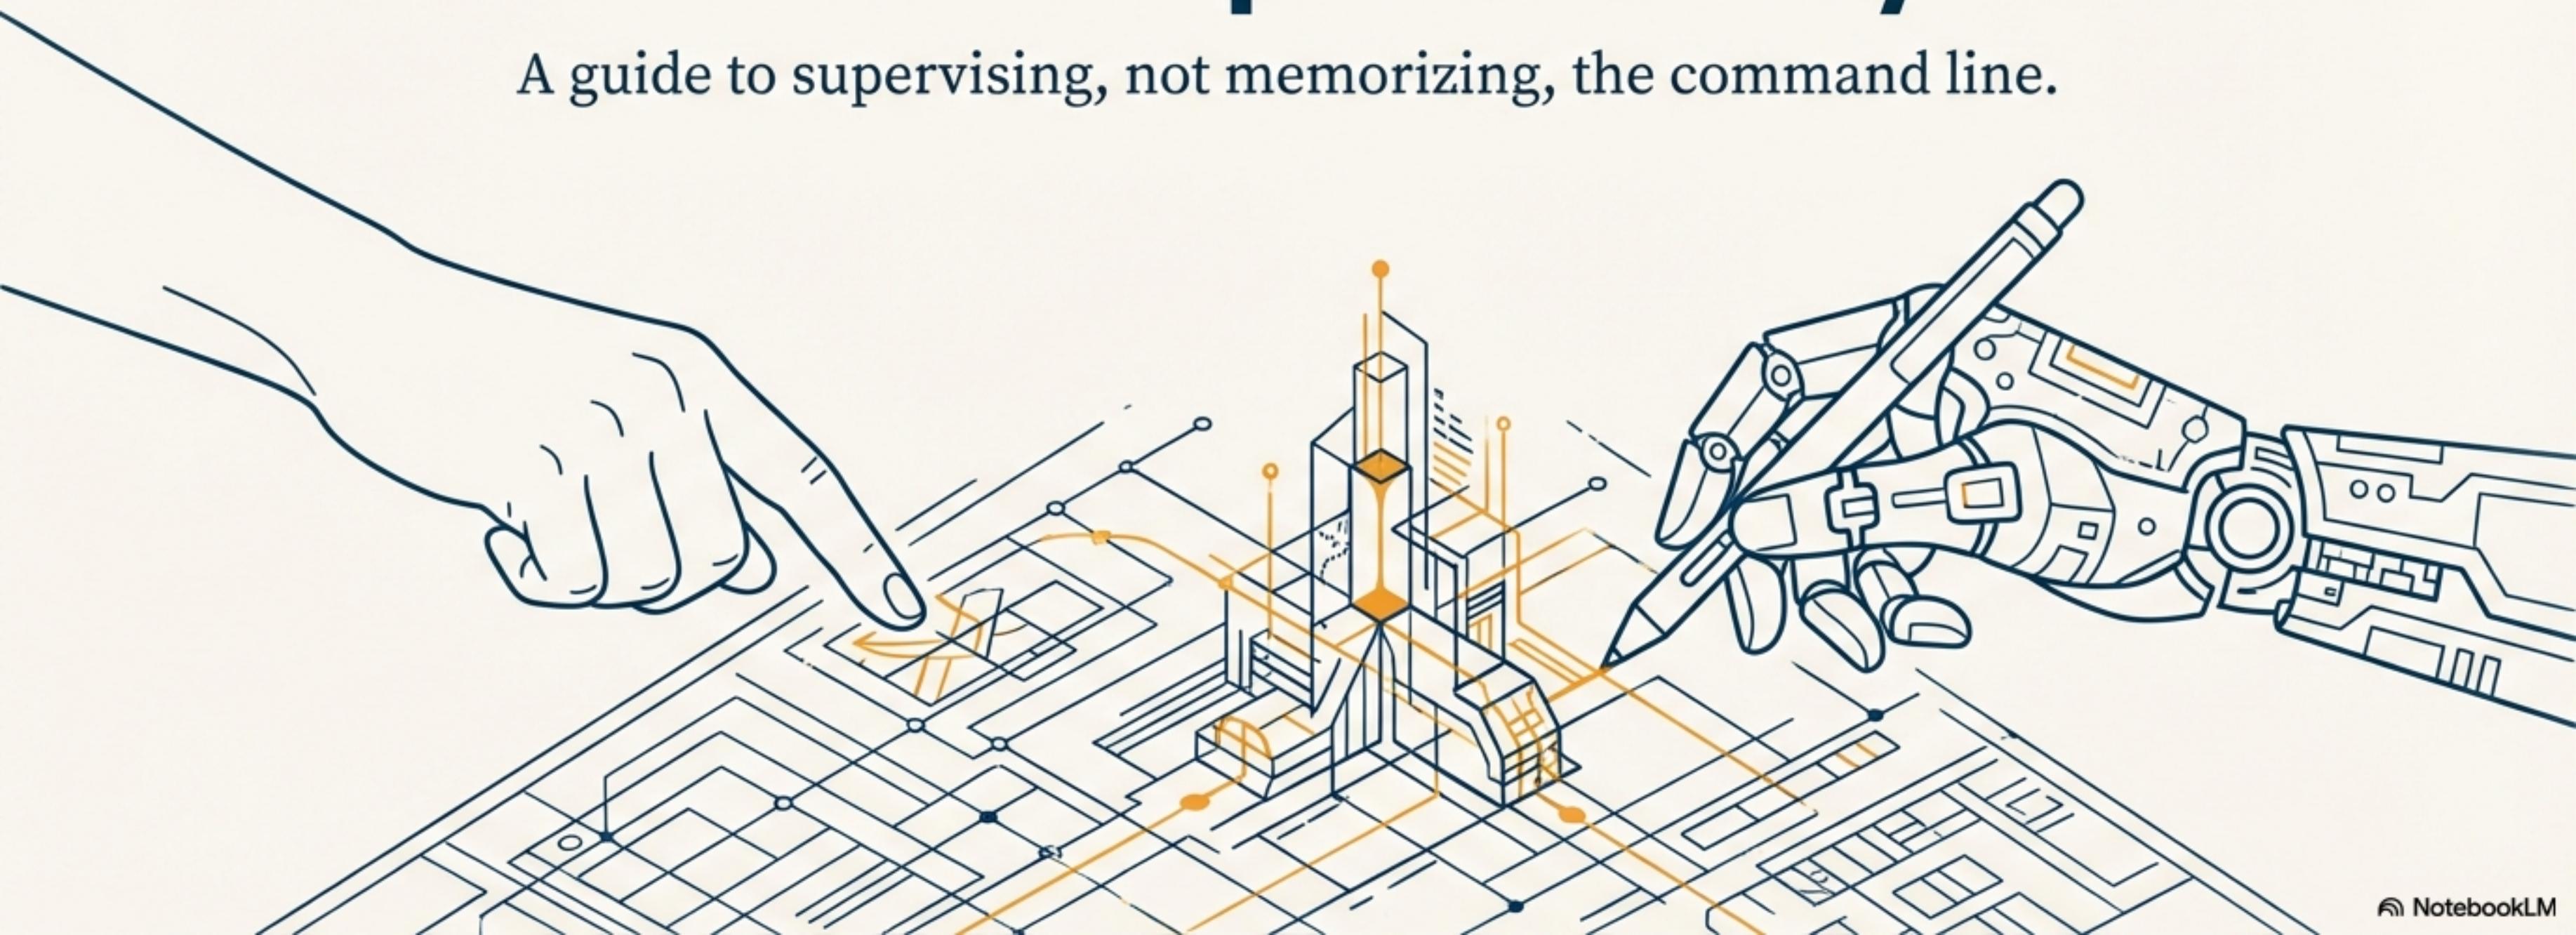


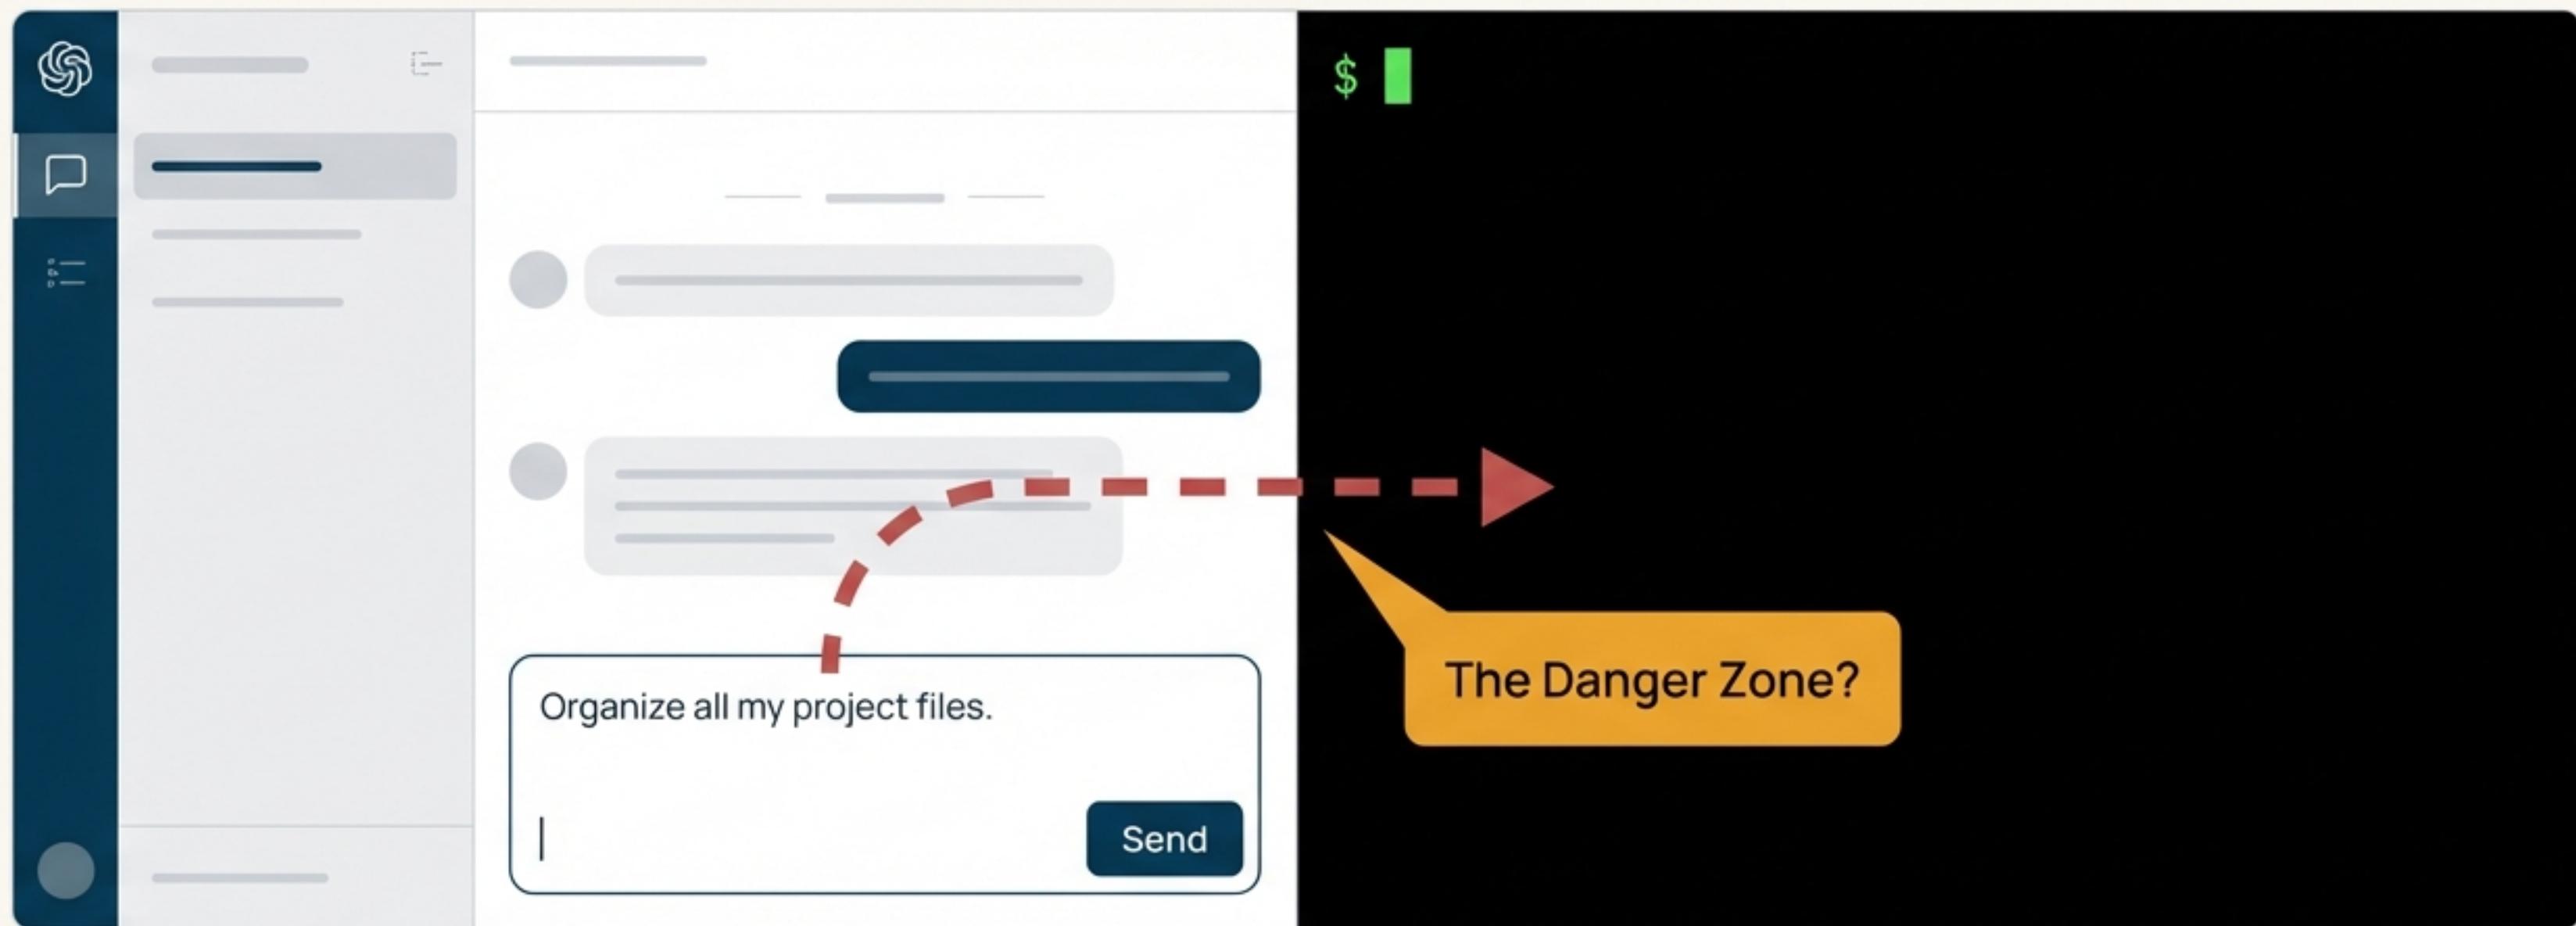
Your AI's Workspace Is Your Responsibility

A guide to supervising, not memorizing, the command line.



You Wouldn't Let a Contractor Work Blind in Your House

Your AI companion is a powerful contractor. When you ask it to “create a project” or “run tests,” it’s swinging a hammer in your computer’s most critical space: the terminal. Without supervision, you’re hoping it doesn’t hit a water pipe.



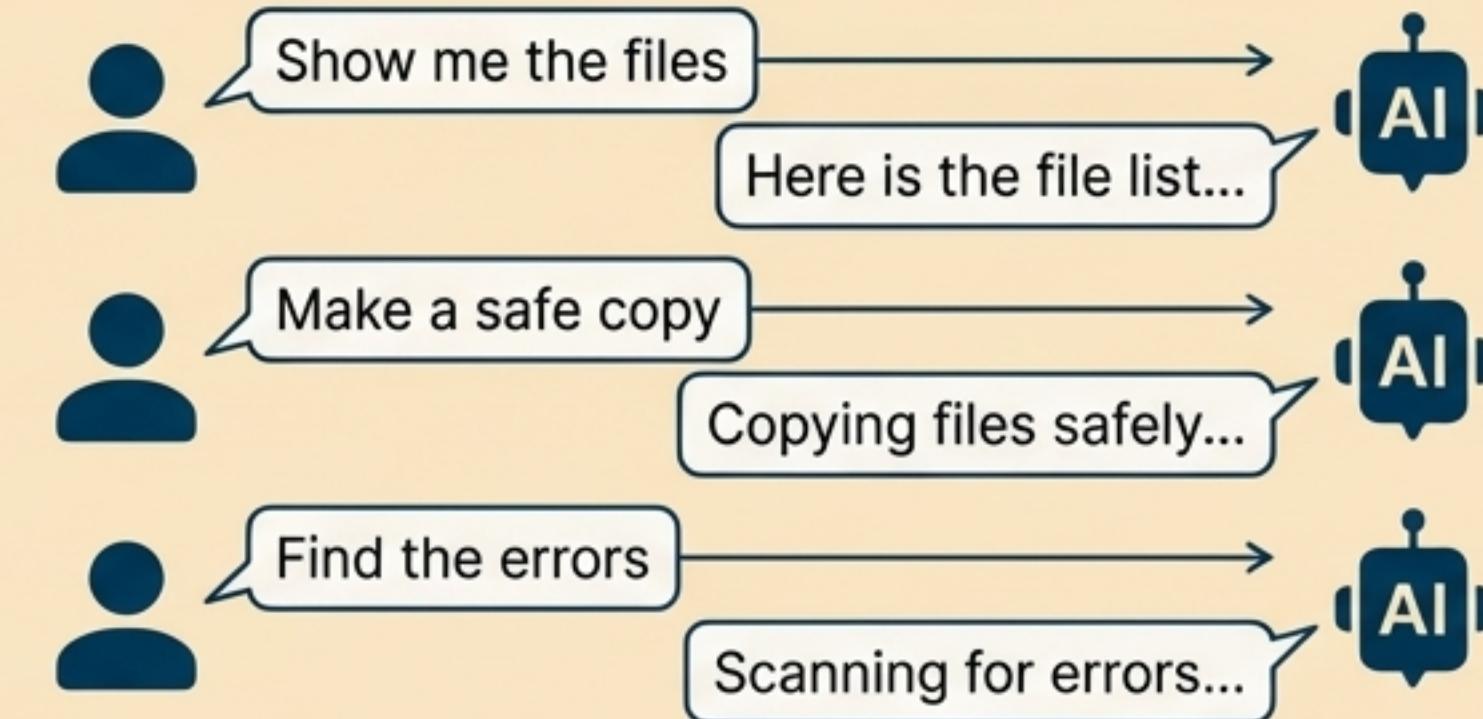
The Goal is Supervision, Not Memorization

You don't need to become a Bash expert. In AI-native development, your role shifts from a command-line operator to a strategic supervisor. Your job is to understand *what* your AI is doing, *why* it matters, and *whether it's safe* to proceed.

OLD WAY

```
ls -lah  
cp -r /path/to/source /path/to/dest  
      grep -r "error" .  
      awk '{print $1}' filename.txt  
      sed 's/foo/bar/g'  
      tar -czvf archive.tar.gz folder  
      ssh user@host chmod 755 script.sh  
      sudo apt-get install package
```

NEW WAY



The Safety-First Dialogue Pattern

Every time you ask your AI to perform an operation, follow this simple, structured dialogue. It's like a surgeon's pre-flight checklist for the command line.



This pattern is your safety net. We will see it applied in every example that follows.

Supervision Starts With One Question: “Where Are You?”

Before any operation, you must know where your AI is “standing” in your file system. The `pwd` (print working directory) command answers this.

You Try It

```
$ pwd  
/Users/yourname/Documents
```

This is your current location.

Your AI Does It



You: Where are you working right now?



AI: I am here: `pwd`

```
/Users/mjs/Documents/code/panaversity-  
official/tutorgpt-build/colearning-python
```

This is the AI's current location.

Your AI isn't using magic. It runs the same commands you do. You can read its output because you understand the tool.

Next, You Ask: “What Can You See Here?”

Once you know the location, you need to see the files and folders there. The `ls -la` command lists everything in detail. Your role is to interpret this evidence.

```
total 24
drwxr-xr-x  4 user  staff   128 Nov  2 10:30 .
drwxr-xr-x 12 user  staff   384 Nov  2 10:25 ..
drwxr-xr-x  7 user  staff   224 Oct 31 15:00 book-source
-rw-r--r--  1 user  staff  1024 Nov  1 09:15 README.md
-rw-r--r--  1 user  staff   450 Oct 28 12:00 config.yaml
```

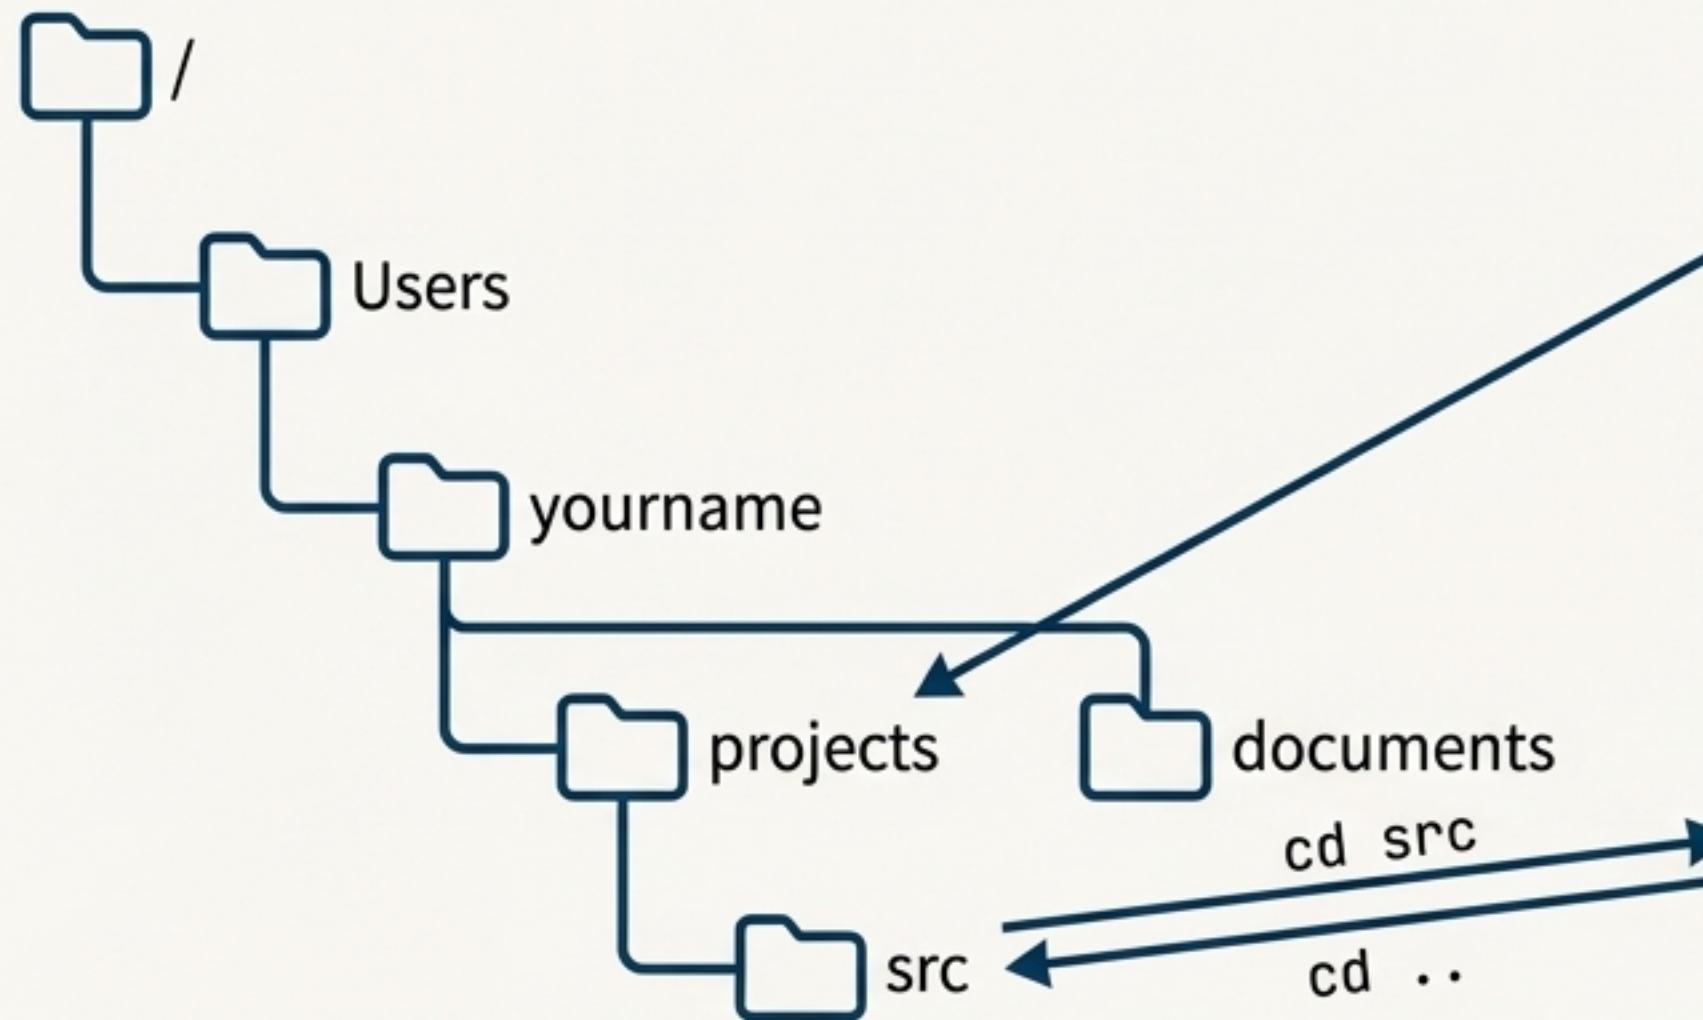
d = directory
(a folder)

- = file
(a regular file)

When your AI says “I can see the `book-source` folder,” you can now look at the `ls` output and verify its claim. You’re reading the evidence yourself.

Supervising Navigation is About Giving Clear Directions

File paths are like addresses. Understanding them is key to preventing your AI from getting lost or operating in the wrong place.



Absolute Paths (A Full Address)

You: "Go directly to my projects folder."

AI: "Okay. cd /Users/yourname/projects"

Starts with '/' or '~'. Works from anywhere.

Relative Paths (Directions from Here)

You: "From here, go into the src folder, then back up."

AI: "Okay. cd src ... Now going up: cd .."

'..' means 'up one level.' Only works from the current location.

An unsafe request is "Delete the backup folder."

A safe request is "Show me which backup folder you see, then delete it."

Not All Operations Are Equal: Know The Risk

Understanding the potential for data loss is central to good supervision. Before approving any file operation, assess its risk level.

Copy

cp

****Low****

Safe—original
is untouched.



Move / Rename

mv

****Medium****

Original location
is gone.



Delete

rm

****High****

Difficult or
impossible to undo.



For any **Medium** or **High** risk operation, the **5-Step Safety Pattern** is non-negotiable.

The Safety Pattern in Action: Preventing Disaster

Let's see what happens when you skip steps versus following the pattern for a high-risk task.

Unsafe: Skipping Steps

 You Delete everything in my `old-projects` folder.

You

 AI Deleting... rm -rf old-projects/ ... Done.

AI



Disaster!

important_data.csv was in that folder and is now gone forever.

Safe: Using the Pattern

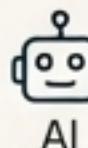
1.  You Delete everything in my `old-projects` folder.

You

2.  AI I'll use rm -rf to delete the folder and everything in it permanently.

4.  You Wait, show me what's in there first.

You

 AI Contents: ls old-projects/`
... project-a, project-b, **important_data.csv**



You Stop! Let me move important_data.csv first.



Safety.

You verified the contents and saved critical data before execution.

Your Code is Public. Your Secrets Are Not.

Hardcoding passwords or API keys in your code is like leaving your house key on the front door. If your code is on GitHub, your secrets are exposed.

Secrets belong in the environment, never in your code.

BAD CODE

```
# DANGEROUS: Do not do this!
api_key = "sk-abc123xyz..."
```

GOOD CODE

```
# SAFE: Read from the environment
import os
api_key = os.getenv("API_KEY")
```

How-To

For temporary use: `export API_KEY="sk-..."` (lasts for one session).

For projects: Store in a `.env` file (and add `.env` to `.gitignore`!).

Package Managers Are Your Software Supplier Network

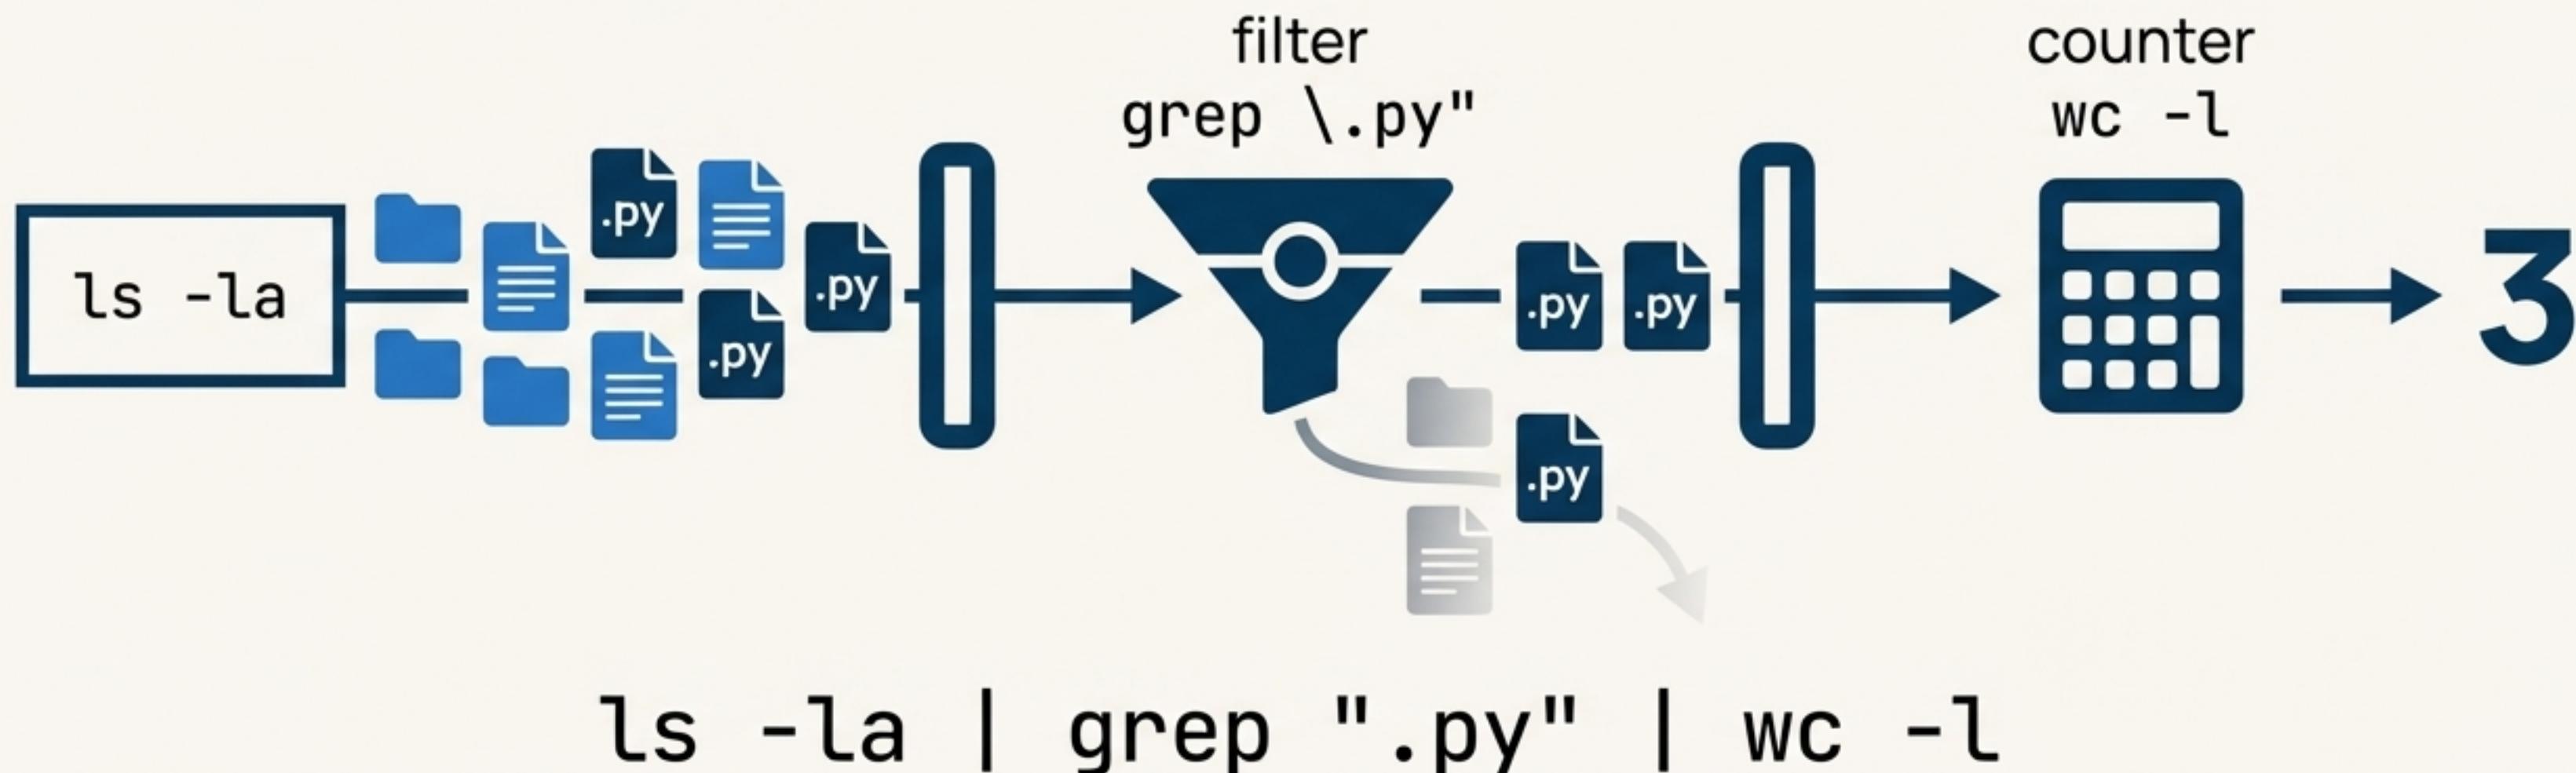
You wouldn't build a bookshelf by chopping down a tree. You order parts. Package managers (`pip` for Python, `npm` for Node.js) do the same for your code. You ask for one thing, and they deliver it along with all the other parts it needs to work (its “dependencies”).



You asked for one package, but five were installed. The system handled the complexity.
Your job is to verify the main package was installed correctly.

Pipes Are an Assembly Line for Your Data

The pipe character `|` chains commands together, sending the output of one command to be the input of the next. This lets you build complex operations from simple, single-purpose tools.



We listed files, filtered for Python files, then counted the results.

You Ask in English. Your AI Builds the Assembly Line.

You don't need to memorize complex commands. Your job is to state your goal clearly and then verify that the AI's proposed data flow makes sense.



> Find all errors in the log file, show me only the unique error messages, and count how many times each one occurred.



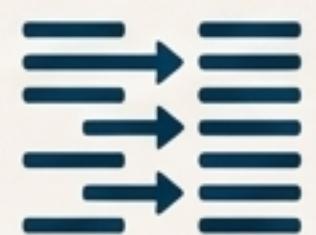
> grep "ERROR" app.log | sort | uniq -c | sort -nr

grep "ERROR"



Finds all
error lines

sort



Groups identical
errors together

uniq -c



Counts each
unique group

sort -nr



Sorts the results to show the
most frequent errors first

Your role is to understand the flow: Find → Group → Count → Sort. The AI handles the syntax.

Let's Review: A Real-World Scenario

Your AI suggests running `rm -rf backup/` to delete old files. Following the Safety-First Pattern, what is the most critical action you must take *before* this command executes?

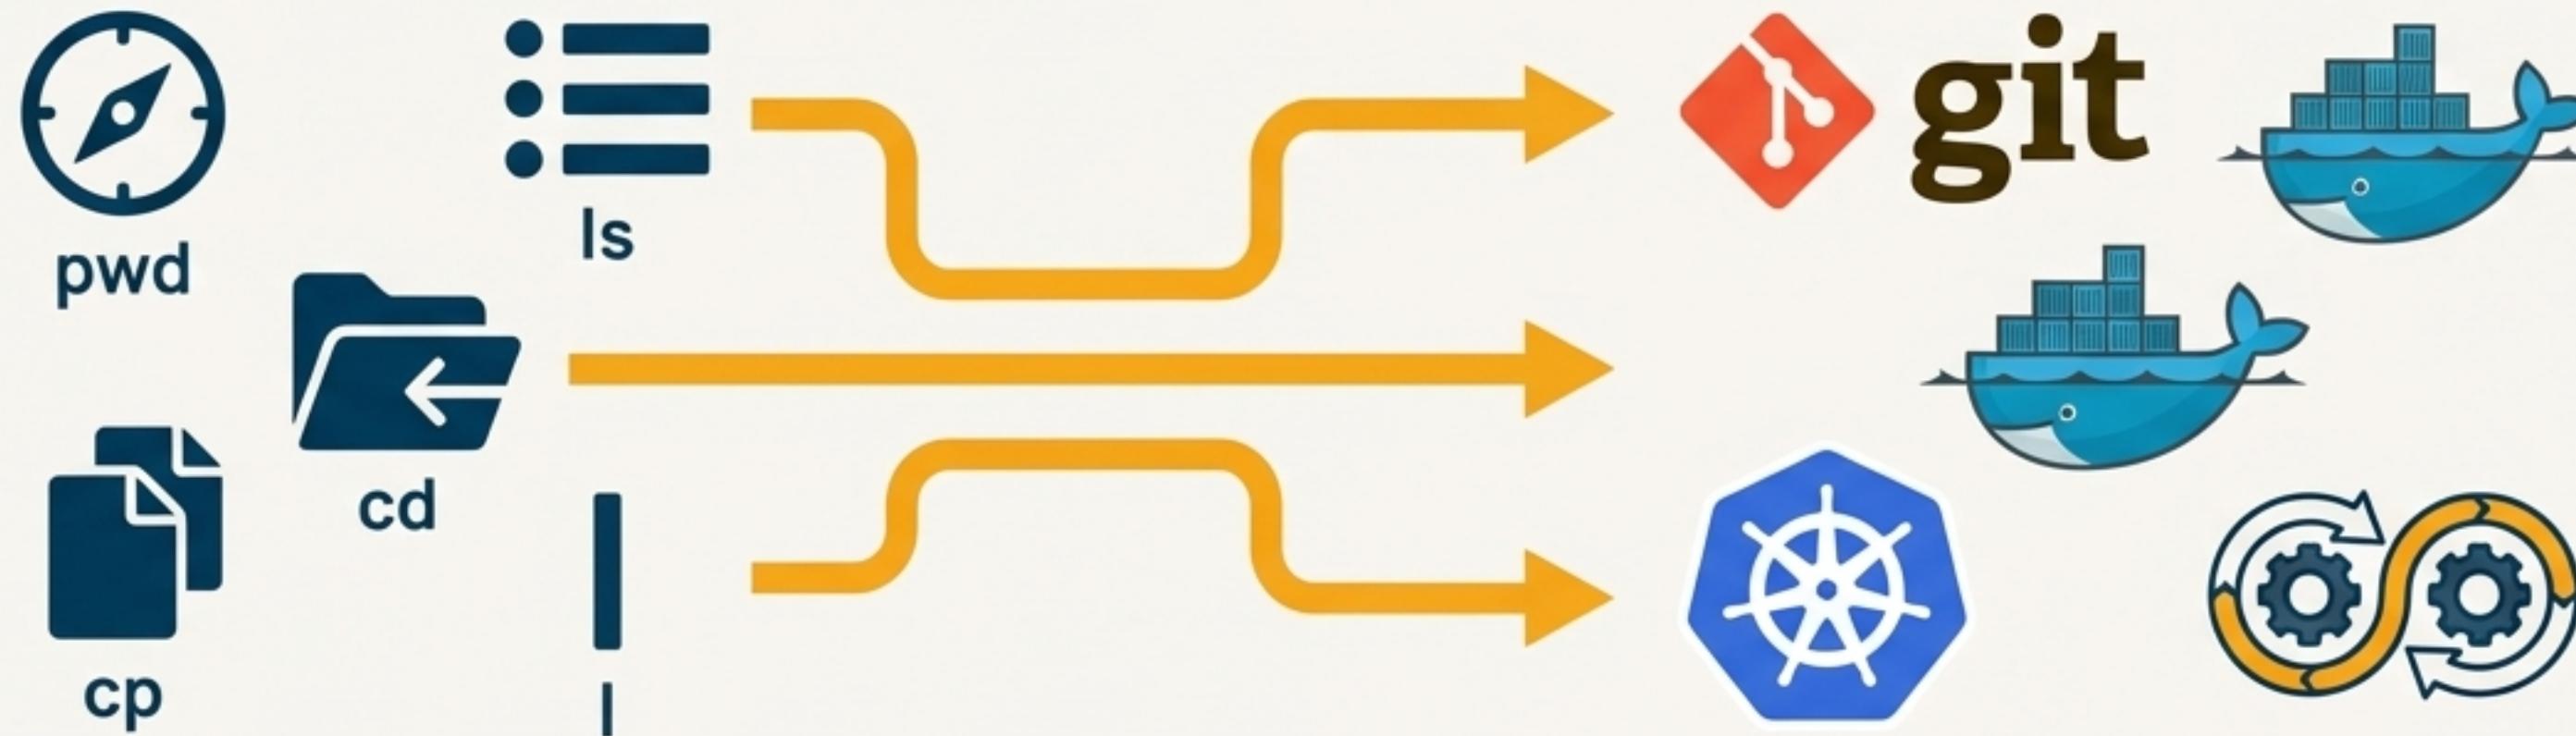
- A) Trust the AI and execute immediately.
- B) First, create a backup of the `backup/` folder.
-  C) **Ask the AI to show its current location (`pwd`) and list the contents of the `backup/` folder (`ls backup/`) for you to verify.**
- D) Memorize what `rm -rf` does from the manual.

Explanation

Correct. Verification before execution is the core of the pattern. You must see WHERE the AI is and WHAT it's about to delete. This prevents accidentally deleting the wrong folder or losing important files. Checking after the fact is too late for a destructive operation.

From Supervisor to Confident Collaborator

Understanding the command line isn't a throwback skill; it's the foundation for all professional development. The safety patterns and verification mindset you've learned here are the same ones you will use to supervise your AI partner across every tool in your stack.



You don't just follow blindly. You collaborate with understanding. This is mastery in the AI era.